

0.1 intro

Our presentation is on a simple proof of sequential work. This paper won the best paper award, and builds upon a paper about **publicly verifiable proofs of sequential work** by Mahmoody, Moran and Vadhan in 2013.

So a proof of sequential work is a protocol for proving that one did sequential computational work related to some statement.

- Ensuring someone did sequential computational work is useful for time-stamping, and the main motivation of this proof is its application in blockchain designs and applicable to cryptocurrencies.
- The proof of sequential work is done with the assumption of a random oracle model. So the prover and verifier, that both have access to a random oracle, which is basically a hash function (will show some properties later)
- so my partner will construct and define a Proof of Sequential Work
- then I will prove two lemmas on the properties of random oracles, 3 graph lemmas and prove the security and efficiency of the protocol

0.2 Lemmas

Lemma 1: Random Oracles are Collision Resistant

- we are given an adversary and a random function. WTS that if A makes at most q queries, the probability that two queries collide is bounded by $\frac{q^2}{2^{w+1}}$
- to prove this, we find the probability of collision between two consecutive queries and add then use the union bound to bound them
- each individual query has length w

Consider an adversary \mathcal{A}^H which is given access to a random function $H : \{0, 1\}^* \rightarrow \{0, 1\}^w$. If \mathcal{A} makes at most q queries, the probability of two colliding queries $x \neq x', H(x) = H(x')$ is at most $\frac{q^2}{2^{w+1}}$

Proof. For individual queries $x_1 \leq x_i \leq x_q$, the probability of collision between the i^{th} query with a previous query, (i.e. $P(H(x_i) = H(x_{i-1}))$) is bounded by $\frac{i-1}{2^w}$. So the probability of collision for all q queries is bounded by $\sum_{i=1}^q \frac{i-1}{2^w} = \frac{(q-1)(q-1+1)}{2 \cdot 2^w} = \frac{q^2}{2^{w+1}}$. \square

Lemma 2: Random Oracles are Sequential

- an H -sequence of length s is a sequence of s strings $x_0, \dots, x_s \in \{0, 1\}^*$, where for each $i < s$, $H(x_i)$ is a substring of x_{i+1} .

Consider an adversary \mathcal{A}^H which is given access to a random function $H : \{0, 1\}^* \rightarrow \{0, 1\}^w$ that it can query for at most $s - 1$ rounds. Each round, \mathcal{A}^H can make arbitrarily many parallel queries. If A makes at most q queries of total length Q bits, then the probability that it outputs an H -sequence x_0, \dots, x_n is at most $q \cdot \frac{Q + \sum_{i=1}^s |x_i|}{2^w}$.

Proof. We can divide this into two cases, where (1) A “gets lucky” with one x_i , or (2) there’s collision, i.e. for some $x_i \neq x_j$, $H(x_i) = H(x_j)$.

- Case 1: for some $0 \leq i < s$, $H(x_i)$ is a substring of x_{i+1} , but A did not query x_i . Since H is a uniformly random function, the probability that $H(x_i) \subseteq_{i+1}$ for some i and some a, b would at most $q \cdot \frac{|x_i|}{2^w}$. Thus the probability for any i is at most $q \cdot \frac{\sum_{i=0}^s |x_i|}{2^w}$, by union bound.
- Case 2: for some $1 \leq i \leq j \leq s - 1$ and some queries x_i, x_j , the probability of collision, i.e. that $x_i \supseteq H(x_j)$ is bounded by $q \cdot \frac{Q}{2^w}$.

Adding the two cases, we get $q \cdot \frac{Q + \sum_{i=1}^s |x_i|}{2^w}$. □

Lemma 3: The labels of G_n^{PoSW} can be computed in topological order using only $w \cdot (n + 1)$ bits of memory

Proof. n is the depth of the graph and w is the output range of the hash function. The proof is a backward induction on the depth of G_n^{PoSW} .

1. First, separate G_n^{PoSW} into Right and Left subtrees. Each subtree is isomorphic to G_{n-1}^{PoSW} , if we don’t take into account the edges going from $label_0$ to leaves on the Right subtree.
2. We calculate $label_0$ on the Left subtree using the space it would take to calculate G_{n-1}^{PoSW} , and keep $label_0$.
3. Then, to calculate $label_1$ on the Right subtree, we need the space it takes to calculate G_{n-1}^{PoSW} , plus w bits to store $label_1$.
4. Then, using only $label_0$ and $label_1$, we calculate the label of the root $label_\epsilon = H(\epsilon, label_0, label_1)$.

Thus, the memory required to compute G_n^{PoSW} is the memory it takes to compute $w + G_{n-1}^{\text{PoSW}} = w + w + G_{n-2}^{\text{PoSW}} = k \cdot w + G_{n-k}^{\text{PoSW}}$.

For base case G_0^{PoSW} , there’s only 1 node, meaning the root can be computed in w bits. So we get $w + G_{n-1}^{\text{PoSW}} = k \cdot w + G_{n-k}^{\text{PoSW}} = n \cdot w + G_0^{\text{PoSW}} = n \cdot w + w = w(n + 1)$. □

Lemma 4: Take a graph $G_n^{\text{PoSW}} = (V, E)$. For any $S \subseteq V$, the subgraph of G_n^{PoSW} consisting of nodes $V - D_{S^*}$ has a directed path going through all the leftover nodes (there are $|V| - |D_{S^*}| = N - |D_{S^*}|$ leftover nodes).

Proof. This proof is an induction on n for G_n^{PoSW} . G_0^{PoSW} is obviously true since it contains a single node. Suppose the lemma holds for G_i^{PoSW} . We now want to show it holds for G_{i+1}^{PoSW} . So pick some $G_{i+1}^{\text{PoSW}} = (V, E)$. G_{i+1}^{PoSW} has a Left and Right subgraph, and root ϵ . The Left and Right subgraphs are isomorphic to G_i^{PoSW} , except for extra edges from label_0 to leaves of the Right subgraph. Consider an arbitrary $S \subseteq V$, and these four cases:

- **case 1:** If node $\epsilon \in S^*$ then we are done because D_{S^*} would be the whole graph and it is vacuously true that $V - D_{S^*}$ has a directed path.
- **case 2:** Suppose nodes $0 \in S^*$, $1 \notin S^*$ then the whole Left subtree would be in D_{S^*} . The Right subtree would become equivalent to G_i^{PoSW} and by assumption the subgraph on $V - D_{S^*}$ has a direct path to 1. Add an edge $1 \rightarrow \epsilon$ and we are done.
- **case 3:** Suppose $0 \notin S^*$, $1 \in S^*$. By the same argument as case 2, we can find a direct path going through the leftover nodes.
- **case 4:** Suppose $0 \notin S^*$, $1 \notin S^*$. Then, take the Left subgraph (equivalent to G_i^{PoSW}) and find a directed path ending in node 0. Take the Right subgraph (equivalent to G_i^{PoSW}) and find a directed path starting at leaf v . Then, link the Left and Right subgraph by adding edges to $0 \rightarrow v$ and $1 \rightarrow \epsilon$.

□

Lemma 5: For any $S^*, S \subset V$, D_{S^*} contains $|\{0, 1\}^n \cap D_{S^*}| = \frac{|D_{S^*}| + |S^*|}{2}$ many leaves

Proof. Suppose $S^* = \{v_1, \dots, v_k\}$. Then, $D_{v_i} \cap D_{v_j} = \emptyset$ for $i \neq j$ because S^* is a minimal set. Thus, to find the total number of leaves in D_{S^*} , we can sum the number of leaves in each D_{v_i} , which is easier since each D_{v_i} is a full binary tree with $\frac{|D_{v_i}| + 1}{2}$ leaves. So

$$\begin{aligned} |\{0, 1\}^n \cap D_{S^*}| &= \sum_{i=1}^k |\{0, 1\}^n \cap D_{v_i}| \\ &= \sum_{i=1}^k \frac{|D_{v_i}| + 1}{2} \\ &= \frac{|D_{S^*}| + |S^*|}{2} \end{aligned}$$

□

0.3 Proof of Security

honest prover will make N queries to prove the statement, and we want to show a dishonest prover must make close to N queries or otherwise get rejected.

Theorem 1: Consider a PoSW defined using parameters N , H , t , and M as defined above, with an additional parameter $\alpha > 0$. α is what the authors call a "soundness gap", which is the percentage difference between N and how many queries to H a cheating prover \tilde{P} actually makes, i.e. a cheating prover \tilde{P} will make at most $(1 - \alpha)N$ queries. For such a PoSW, the verifier \mathcal{V} will reject with probability $1 - (1 - \alpha)^t - \frac{2 \cdot n \cdot w \cdot q^2}{2^w}$.

the first guy (1-a) represents the answer the cheater gave is within soundness gap

the second guy (2nwq...) is the event where there exists a collision or P breaks sequentiality so the cheater gets away

- So $\frac{2 \cdot n \cdot w \cdot q^2}{2^w}$ accounts for the assumptions we make bounded by lemma 1 and lemma 2, that there will be a collision and that P will break sequentiality of H : $\frac{q^2}{2^{w+1}} + q \frac{Q + \sum_i^s |x_i|}{2^w}$
- Q is the total number of bits queries, $Q \leq q \cdot w \cdot (n + 1)$, where $w \cdot (n + 1)$ is the largest possible query. And q is the total number of queries.
- $|x_i|$ in this case is the sequence of labels which is bounded by w bits, and s is the number of rounds queried and is bounded by q (the total number of queries).
- so $\leq q \frac{q \cdot w \cdot (n+1) + qw}{2^w} + \frac{q^2/2}{2^w} = \frac{q^2 w(n+1) + q^2 w + q^2/2}{2^w} = \frac{q^2(w(n+1) + w + 1/2)}{2^w} < \frac{q^2(2wn)}{2^w}$

Proof: Consider the probability that \mathcal{V} detects an inconsistent answer to the challenge (challenge is a set of nodes and the answer should be their corresponding labels).

- Let set $S \subseteq V = \{0, 1\}^{\leq n}$ be the set of inconsistent vertices. By Lemma 4 there is a path going through all the vertices of $V - D_{S^*}$ (which are all labeled correctly)
- so the path is an H_χ -sequence of the leftover nodes $N - |D_{S^*}|$. (Where N is the number of all nodes.)

Now we can divide this into two cases:

- the case where the answer is within the soundness gap, then the verifier considers that close enough to N and won't reject it
- the answer is not within the soundness gap, so the verifier will reject with high probability

Case 1 ($|D_{S^*}| \leq \alpha N$): $\tilde{\mathcal{P}}$ must have made at least $(1 - \alpha)N$ sequential queries to H_χ , to compute the H_χ -sequence of length $N - |D_{S^*}|$. verifier accepts.

Case 2 ($|D_{S^*}| > \alpha N$): By definition, the number of nodes is $N = (2^{n+1} - 1)$, so $\alpha N = \alpha(2^{n+1} - 1)$.

By Lemma 5, D_{S^*} contains $\frac{|D_{S^*}| + |S^*|}{2}$ leaves. Substituting this into $|D_{S^*}| > \alpha N$, and looking at the leaves $\alpha N = \alpha 2^{n+1} - \alpha \frac{|D_{S^*}| + |S^*|}{2} > \frac{\alpha(2^{n+1} + |S^*|)}{2} > \frac{\alpha 2^{n+1}}{2} = \alpha 2^n$ we get the number of leaves $|\{0, 1\}^n \cap D_{S^*}| = \frac{|D_{S^*}| + |S^*|}{2} > \alpha 2^n$.

So the case is not within soundness gap.

So the general idea, \mathcal{V} will reject τ (the answer to the challenge γ_i , $\gamma = (\gamma_1, \dots, \gamma_t)$), if there exists a node $u \in S$ that's on the path from γ_i to the root, i.e.

\mathcal{V} rejects if $\gamma \cap D_{S^*} = \gamma \cap \hat{S}^* = \gamma \cap \hat{S} \neq \emptyset$.

From the previous inequality on the number of leaves and assuming that all γ_i 's are sampled uniformly, we get $\Pr[\gamma_i \notin D_{S^*}] = 1 - |\{0, 1\}^n \cap D_{S^*}| / 2^n < 1 - \alpha$.

(probability of challenged nodes aren't in the path of inconsistent answers)
And since all γ_i 's are sampled independently, $\Pr[\gamma \cap D_{S^*} = \emptyset] = \prod_{i=1}^t \Pr[\gamma_i \notin D_{S^*}] < (1 - \alpha)^t$.

Combined all together, a cheating prover $\tilde{\mathcal{P}}$ will have its proof rejected with probability $1 - (1 - \alpha)^t - \frac{2 \cdot n \cdot w \cdot q^2}{2^w}$. 1-(inconsistent leaves)

0.4 Proof of Efficiency

0.4.1 proof size

w bits specify a label and n bits specify a node. Thus, the exchanged messages and their lengths are as follows:

- $|\chi| = w$. χ is the initial statement that is a uniformly random w -bit string. It is initially communicated from Verifier \rightarrow Prover.

- $|\phi| = w$. ϕ and ϕ_p are proofs computed from PoSW. ϕ is the root label sent from Prover \rightarrow Verifier. It is a w -bit string, since labels are calculated using $H : \{0, 1\}^{\leq w(n+1)} \rightarrow \{0, 1\}^w$.
- $|\gamma| = t \cdot n$. $\gamma = (\gamma_1, \dots, \gamma_t)$ is a challenge sent from Verifier \rightarrow Prover. It consists of t leaf nodes of n bits.
- $|\tau| \leq t \cdot w \cdot n$. $\tau := \text{open}(\chi, N, \phi_p, \gamma)$ is the answer sent from Prover \rightarrow Verifier, which answers the challenge γ . The answer will contain the w -bit label for each n -bit γ_i .

0.4.2 prover efficiency

The prover \mathcal{P} 's efficiency depends queries made while computing the PoSW and open.

- $\text{PoSW}^{H_x}(N)$ is computed using N sequential queries to H_x . Each input has a length of at most $(n+1) \cdot w$ bits, by definition.
- $\text{open}^{H_x}(N, \phi, \gamma) = \tau$. open requires
 1. $(n+1)w$ bits to compute each label of the challenge,
 2. $2^{m+1}w$ labels to be stored in ϕ_p , and
 3. $|\tau| \leq t \cdot w \cdot n$ bits to send back.

Adding these, we need $(n+1+n \cdot t + 2^{m+1})w$ bits of memory. We examine the different cases depending on m , i.e. how many levels are used to store ϕ_p :

Case $m = n$ \mathcal{P} stores all the labels computed by $\text{PoSW}^{H_x}(N)$, so no additional queries are needed

Case $m = 0$, \mathcal{P} does not store any label computed by $\text{PoSW}^{H_x}(N)$, and needs to recompute all N queries

Case $0 < m < n$ Since \mathcal{P} stored the top m levels, it needs to recalculate any query between level n to m . This would require calculating the leaves starting from the $m - n^{th}$ level which would require $(2^{m-n+1} - 1) \cdot t$ queries, for t challenges.

0.4.3 verifier efficiency

The verifier only needs to sample a random challenge of $|\gamma| = t \cdot n$, and computing $\text{verify}(\chi, N, \phi, \gamma, \tau)$. verify makes $t \cdot n$ queries (for each γ) each of length $n \cdot w$ bits (n leaf nodes' length and w label lengths).

and we want to know if A makes at most q queries what's the probability of two queries to collide

1 previous papers

MMV's paper defined their proof of sequential work mainly in the context of time-lock puzzles

it's also possible to use modular exponentiation as a proof of sequential work in the context of CPU benchmarks and time-lock puzzles